

Linear and Neural Sentiment Classification

Qifan Wen

wen.679@osu.edu

1 Logistic Regression

Parameter	Value
Initial learning rate	0.5
Learning rate decay	0.95 per epoch
Batch size	1 (SGD)
Number of epochs	30
Weight initialization	$U(-0.1, 0.1)$
Numerical stability	Clip to $[-500, 500]$

Table 1: Logistic Regression training configuration.
All models use seed 42.




Figure 1: Training loss across all 5 LR configurations.

	Unigram	Bigram	Better
Train Accuracy	99.9%	100.0%	100.0%
Dev Accuracy	77.5%	77.2%	77.1%
Precision	77.1%	76.3%	76.1%
Recall	79.5%	80.0%	80.2%
F1 Score	78.3%	78.1%	78.1%

Table 2: Feature extractor comparison ($\text{lr}=0.5$, $\text{decay}=0.95$).




Figure 2: Training (solid) and dev (dashed) accuracy across all 5 LR configurations.

The improved logistic regression models incorporate three enhancements: (1) **L2 regularization** ($\lambda = 0.1$) prevents overfitting by penalizing large weights, (2) **TF-IDF weighting** emphasizes discriminative words over common ones, and (3) **feature frequency thresholding** ($\text{min_count}=2$) removes rare/noisy features.

As shown in Figure ??, the training loss curves are smoother with L2 regularization, and aggressive decay (0.8) still plateaus at higher loss. Figure ?? shows that L2 regularization prevents perfect training accuracy, which is desirable for generalization, while the dev accuracy curves demon-

	Fixed (1.0)	Default (0.95)	Aggressive (0.8)
Train Acc	100.0%	99.9%	99.4%
Dev Acc	77.8%	77.5%	78.0%
Precision	77.2%	77.1%	77.2%
Recall	80.0%	79.5%	80.6%
F1 Score	78.5%	78.3%	78.9%

Table 3: Learning rate schedule comparison (Unigram features).

strate more stable convergence across all configurations. TF-IDF weighting helps the Unigram model by down-weighting common words like “the” and “a” while emphasizing sentiment-bearing terms. Feature thresholding reduces noise from rare bigrams/trigrams that appear only once in training.

2 Deep Averaging Network

Parameter	Value
Learning Rate	0.0005
Batch Size	1 or 32
Epochs	20
Hidden Size	150
Dropout	0.3
Weight Decay	1e-5
Optimizer	Adam

Table 4: Neural model training configuration. All models use frozen 300d GloVe embeddings, NLLoss, and Xavier initialization with seed 42.

Model	Train	Dev	Prec	Rec	F1
DAN	89.1%	76.7%	78.1%	75.5%	76.7%
DAN+B	92.1%	78.3%	76.7%	82.4%	79.5%
LSTM	99.4%	78.9%	75.9%	85.8%	80.5%
CNN	100.0%	82.6%	82.4%	83.6%	83.0%

Table 5: Neural model performance comparison. DAN+B denotes DAN with mini-batch training (batch size 32).




Figure 4: Training (solid) and dev (dashed) accuracy for all 4 neural models over 20 epochs.

The CNN achieves the best dev accuracy (82.6%) and F1 score (83.0%), demonstrating that local n-gram patterns captured by convolutional filters are highly effective for sentiment classification. Mini-batch training improves DAN generalization (+1.6% dev accuracy) by providing regularization through noisier gradient estimates. LSTM achieves near-perfect training accuracy (99.4%) but shows more variance in dev accuracy, indicating some sensitivity to sequential patterns.

As shown in Figure ??, CNN and LSTM converge to much lower training loss compared to DAN models, correlating with their higher model capacity. The accuracy curves (Figure ??) show CNN reaching 100% training accuracy fastest, while DAN models plateau around 90%; on the dev set, CNN maintains the most stable generalization, peaking around 82% dev accuracy. All neural models outperform the LR baseline (77.5% dev), with frozen GloVe embeddings and weight decay regularization providing a strong foundation. The precision-recall trade-off varies: CNN achieves the best balance (82.4%/83.6%), while LSTM favors recall (85.8%) over precision (75.9%).




Figure 3: Training loss for all 4 neural models over 20 epochs.