

# Linear and Neural Sentiment Classification

Qifan Wen

wen.679@osu.edu

## 1 Logistic Regression

### 1.1 Training Configuration

Parameter	Value
Initial learning rate	0.5
Learning rate decay	0.95 per epoch
Batch size	1 (SGD)
Number of epochs	30
Weight initialization	$U(-0.1, 0.1)$
Random seed	42
Numerical stability	Clip to $[-500, 500]$

Table 1: Logistic Regression training configuration.

### 1.2 Results Table

	Unigram	Bigram	Better
Train Accuracy	99.9%	100.0%	100.0%
Dev Accuracy	77.5%	77.2%	77.1%
Precision	77.1%	76.3%	76.1%
Recall	79.5%	80.0%	80.2%
F1 Score	78.3%	78.1%	78.1%

Table 2: Feature extractor comparison ( $lr=0.5$ , decay=0.95).

	Fixed (1.0)	Default (0.95)	Aggressive (0.8)
Train Acc	100.0%	99.9%	99.4%
Dev Acc	77.8%	77.5%	78.0%
Precision	77.2%	77.1%	77.2%
Recall	80.0%	79.5%	80.6%
F1 Score	78.5%	78.3%	78.9%

Table 3: Learning rate schedule comparison (Unigram features).

### 1.3 Loss Over Epochs Plot

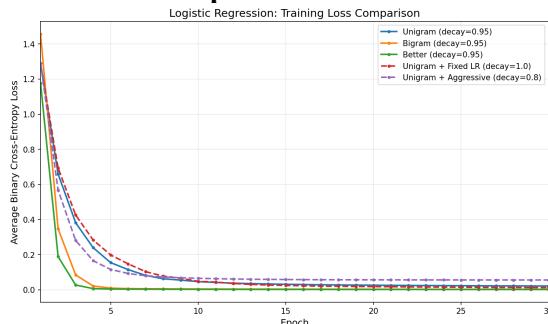


Figure 1: Training loss comparison across all 5 LR configurations. Richer features (Bigram, Better) converge fastest; aggressive decay plateaus early at higher loss.

### 1.4 Analysis

All three feature extractors achieve similar dev accuracy (77.1%–77.5%), with Unigram performing best despite being simplest. Bigram and Better achieve perfect training accuracy (100%) but do not improve generalization, suggesting overfitting—adding more features allows memorization without better discriminative power. Aggressive learning rate decay (0.8) yields the best dev accuracy (78.0%) and F1 (78.9%) despite lowest train accuracy (99.4%), acting as implicit regularization.

As shown in Figure ??, richer features (Bigram, Better) converge fastest (by epoch 4–5) due to larger feature spaces (86K–192K vs 15K for Unigram), while aggressive decay plateaus early ( $\sim 0.056$  final loss) as the learning rate diminishes too quickly. All models show overfitting (train  $\sim 100\%$  vs dev  $\sim 77\%$ –78%), with recall consistently higher than precision across configurations.

## 2 Deep Averaging Network

### 2.1 Training Configuration

Parameter	DAN	DAN+B	LSTM	CNN
Learning Rate	0.001	0.001	0.001	0.001
Batch Size	1	32	32	32
Epochs	10	10	10	10
Hidden Size	100	100	100	100
Dropout	0.3	0.3	0.3	0.3
Optimizer	Adam	Adam	Adam	Adam

Table 4: Neural model training configuration. All models use frozen 300d GloVe embeddings, NLLLoss, and Xavier initialization with seed 42.

### Architecture Notes:

- **DAN:** Avg embedding  $\rightarrow$  FC(300  $\rightarrow$  100)  $\rightarrow$  ReLU  $\rightarrow$  Drop  $\rightarrow$  FC(100  $\rightarrow$  100)  $\rightarrow$  ReLU  $\rightarrow$  Drop  $\rightarrow$  FC(100  $\rightarrow$  2)  $\rightarrow$  LogSoftmax
- **LSTM:** BiLSTM(300  $\rightarrow$  100)  $\rightarrow$  Concat hidden  $\rightarrow$  Drop  $\rightarrow$  FC(200  $\rightarrow$  2)  $\rightarrow$  LogSoftmax

- **CNN:** Conv1d (kernels  $2,3,4,5 \times 25$ )  $\rightarrow$  ReLU  $\rightarrow$  MaxPool  $\rightarrow$  Concat  $\rightarrow$  Drop  $\rightarrow$  FC(100  $\rightarrow$  2)  $\rightarrow$  LogSoftmax

## 2.2 Results Table

Model	Train	Dev	Prec	Rec	F1
DAN	84.1%	77.4%	77.1%	79.1%	78.1%
DAN+B	85.3%	78.6%	75.1%	86.5%	80.4%
LSTM	98.1%	82.5%	79.9%	87.6%	83.6%
CNN	99.8%	81.2%	79.9%	84.2%	82.0%

Table 5: Neural model performance comparison. DAN+B denotes DAN with mini-batch training (batch size 32).

## 2.3 Loss Over Epochs Plot

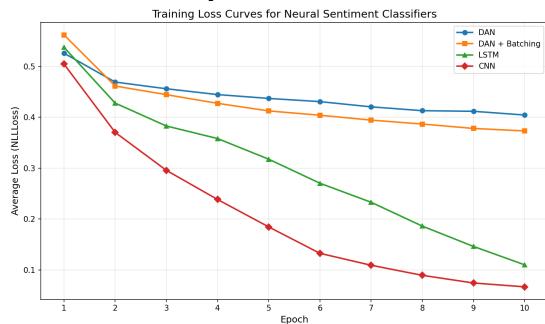


Figure 2: Training loss curves for all 4 neural models over 10 epochs. DAN models plateau at higher loss ( $\sim 0.37\text{--}0.40$ ), while LSTM and CNN achieve much lower final loss ( $\sim 0.07\text{--}0.11$ ), correlating with their higher train accuracies.

## 2.4 Analysis

The LSTM achieves the best dev accuracy (82.5%) and F1 score (83.6%), demonstrating that sequential modeling captures valuable word order information that averaging-based DAN approaches miss. Mini-batch training improves DAN generalization (+1.2% dev accuracy, +2.3% F1) by providing regularization through noisier gradient estimates. CNN achieves strong train accuracy (99.8%) but slightly lower dev accuracy (81.2%) than LSTM, suggesting overfitting to local n-gram patterns rather than global sentence semantics.

As shown in Figure ??, CNN and LSTM converge to much lower training loss ( $\sim 0.07\text{--}0.11$ ) compared to DAN models ( $\sim 0.37\text{--}0.40$ ), correlating with their higher model capacity and train accuracies. All neural models outperform the LR baseline (77.5% dev), with frozen GloVe embeddings providing a strong foundation. The precision-recall trade-off varies across models: LSTM achieves the best balance (79.9%/87.6%), while DAN+Batching favors recall (86.5%) over precision (75.1%).