

Linear and Neural Sentiment Classification

Qifan Wen

wen.679@osu.edu

1 Logistic Regression

1.1 Training Configuration

| Parameter | Value |
|-----------------------|-----------------------|
| Initial learning rate | 0.5 |
| Learning rate decay | 0.95 per epoch |
| Batch size | 1 (SGD) |
| Number of epochs | 30 |
| Weight initialization | $U(-0.1, 0.1)$ |
| Random seed | 42 |
| Numerical stability | Clip to $[-500, 500]$ |

Table 1: Logistic Regression training configuration.

1.2 Results Table

| | Unigram | Bigram | Better |
|----------------|---------|--------|--------|
| Train Accuracy | 99.9% | 100.0% | 100.0% |
| Dev Accuracy | 77.5% | 77.2% | 77.1% |
| Precision | 77.1% | 76.3% | 76.1% |
| Recall | 79.5% | 80.0% | 80.2% |
| F1 Score | 78.3% | 78.1% | 78.1% |

Table 2: Feature extractor comparison ($lr=0.5$, $decay=0.95$).

| | Fixed (1.0) | Default (0.95) | Aggressive (0.8) |
|-----------|----------------|-------------------|---------------------|
| Train Acc | 100.0% | 99.9% | 99.4% |
| Dev Acc | 77.8% | 77.5% | 78.0% |
| Precision | 77.2% | 77.1% | 77.2% |
| Recall | 80.0% | 79.5% | 80.6% |
| F1 Score | 78.5% | 78.3% | 78.9% |

Table 3: Learning rate schedule comparison (Unigram features).

1.3 Loss Over Epochs Plot

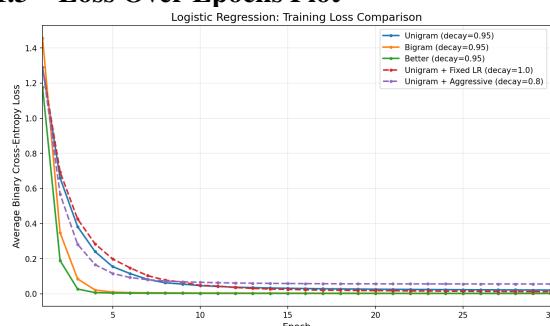


Figure 1: Training loss comparison across all 5 LR configurations. Richer features (Bigram, Better) converge fastest; aggressive decay plateaus early at higher loss.

1.4 Analysis

All three feature extractors achieve similar dev accuracy (77.1%–77.5%), with Unigram performing best despite being simplest. Bigram and Better achieve perfect training accuracy (100%) but do not improve generalization, suggesting overfitting—adding more features allows memorization without better discriminative power. Aggressive learning rate decay (0.8) yields the best dev accuracy (78.0%) and F1 (78.9%) despite lowest train accuracy (99.4%), acting as implicit regularization.

As shown in Figure 1, richer features (Bigram, Better) converge fastest (by epoch 4–5) due to larger feature spaces (86K–192K vs 15K for Unigram), while aggressive decay plateaus early (~ 0.056 final loss) as the learning rate diminishes too quickly. All models show overfitting (train $\sim 100\%$ vs dev $\sim 77\text{--}78\%$), with recall consistently higher than precision across configurations.

2 Deep Averaging Network

2.1 Training Configuration

| Parameter | DAN | DAN+B | LSTM | CNN |
|---------------|-------|-------|-------|-------|
| Learning Rate | 0.001 | 0.001 | 0.001 | 0.001 |
| Batch Size | 1 | 32 | 32 | 32 |
| Epochs | 10 | 10 | 10 | 10 |
| Hidden Size | 100 | 100 | 100 | 100 |
| Dropout | 0.3 | 0.3 | 0.3 | 0.3 |
| Optimizer | Adam | Adam | Adam | Adam |

Table 4: Neural model training configuration. All models use frozen 300d GloVe embeddings, NLLLoss, and Xavier initialization with seed 42.

Architecture Notes:

- **DAN:** Avg embedding \rightarrow FC(300 \rightarrow 100) \rightarrow ReLU \rightarrow Drop \rightarrow FC(100 \rightarrow 100) \rightarrow ReLU \rightarrow Drop \rightarrow FC(100 \rightarrow 2) \rightarrow LogSoftmax
- **LSTM:** BiLSTM(300 \rightarrow 100) \rightarrow Concat hidden \rightarrow Drop \rightarrow FC(200 \rightarrow 2) \rightarrow LogSoftmax
- **CNN:** Conv1d (kernels 2,3,4,5 \times 25) \rightarrow

ReLU → MaxPool → Concat → Drop →
FC(100→2) → LogSoftmax

2.2 Results Table

| Model | Train | Dev | Prec | Rec | F1 |
|-------|-------|-------|-------|-------|-------|
| DAN | 84.1% | 77.4% | 77.1% | 79.1% | 78.1% |
| DAN+B | 85.3% | 78.6% | 75.1% | 86.5% | 80.4% |
| LSTM | 98.1% | 82.5% | 79.9% | 87.6% | 83.6% |
| CNN | 99.8% | 81.2% | 79.9% | 84.2% | 82.0% |

Table 5: Neural model performance comparison. DAN+B denotes DAN with mini-batch training (batch size 32).

2.3 Loss Over Epochs Plot

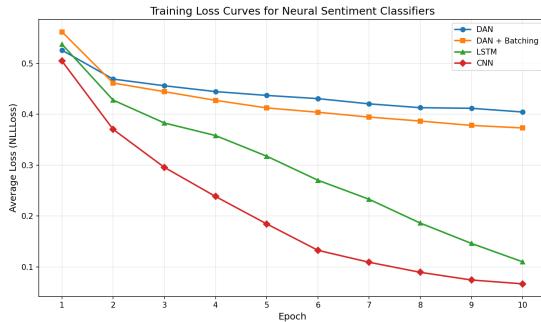


Figure 2: Training loss curves for all 4 neural models over 10 epochs. DAN models plateau at higher loss ($\sim 0.37\text{--}0.40$), while LSTM and CNN achieve much lower final loss ($\sim 0.07\text{--}0.11$), correlating with their higher train accuracies.

2.4 Analysis

The LSTM achieves the best dev accuracy (82.5%) and F1 score (83.6%), demonstrating that sequential modeling captures valuable word order information that averaging-based DAN approaches miss. Mini-batch training improves DAN generalization (+1.2% dev accuracy, +2.3% F1) by providing regularization through noisier gradient estimates. CNN achieves strong train accuracy (99.8%) but slightly lower dev accuracy (81.2%) than LSTM, suggesting overfitting to local n-gram patterns rather than global sentence semantics.

As shown in Figure 2, CNN and LSTM converge to much lower training loss ($\sim 0.07\text{--}0.11$) compared to DAN models ($\sim 0.37\text{--}0.40$), correlating with their higher model capacity and train accuracies. All neural models outperform the LR baseline (77.5% dev), with frozen GloVe embeddings providing a strong foundation. The precision-recall trade-off varies across models: LSTM achieves the best balance (79.9%/87.6%), while DAN+Batching favors recall (86.5%) over precision (75.1%).